

# Q-Learning

**Shusen Wang**

<http://wangshusen.github.io/>

# Sarsa VS Q-Learning

- Sarsa is for training action-value function,  $Q_{\pi}(s, a)$ .
- TD target:  $y_t = r_t + \gamma \cdot Q_{\pi}(s_{t+1}, a_{t+1})$ .
- We used Sarsa for updating value network (critic).

# Sarsa VS Q-Learning

- Q-learning is for training the optimal action-value function,  $Q^*(s, a)$ .
- TD target:  $y_t = r_t + \gamma \cdot \max_a Q^*(s_{t+1}, a)$ .
- We used Q-learning for updating DQN.

**Derive TD Target**

# Derive TD Target

- We have proved that for all  $\pi$ ,

$$Q_{\pi}(s_t, a_t) = \mathbb{E}[R_t + \gamma \cdot Q_{\pi}(S_{t+1}, A_{t+1})].$$

# Derive TD Target

- We have proved that for all  $\pi$ ,

$$Q_{\pi}(s_t, a_t) = \mathbb{E}[R_t + \gamma \cdot Q_{\pi}(S_{t+1}, A_{t+1})].$$

- If  $\pi$  is the optimal policy  $\pi^*$ , then

$$Q_{\pi^*}(s_t, a_t) = \mathbb{E}[R_t + \gamma \cdot Q_{\pi^*}(S_{t+1}, A_{t+1})].$$

# Derive TD Target

- We have proved that for all  $\pi$ ,

$$Q_{\pi}(s_t, a_t) = \mathbb{E}[R_t + \gamma \cdot Q_{\pi}(S_{t+1}, A_{t+1})].$$

- If  $\pi$  is the optimal policy  $\pi^*$ , then

$$Q_{\pi^*}(s_t, a_t) = \mathbb{E}[R_t + \gamma \cdot Q_{\pi^*}(S_{t+1}, A_{t+1})].$$

- $Q_{\pi^*}$  and  $Q^*$  both denote *the optimal action-value function*.

# Derive TD Target

- We have proved that for all  $\pi$ ,

$$Q_{\pi}(s_t, a_t) = \mathbb{E}[R_t + \gamma \cdot Q_{\pi}(S_{t+1}, A_{t+1})].$$

- If  $\pi$  is the optimal policy  $\pi^*$ , then

$$Q_{\pi^*}(s_t, a_t) = \mathbb{E}[R_t + \gamma \cdot Q_{\pi^*}(S_{t+1}, A_{t+1})].$$

- $Q_{\pi^*}$  and  $Q^*$  both denote *the optimal action-value function*.

**Identity:**  $Q^*(s_t, a_t) = \mathbb{E}[R_t + \gamma \cdot Q^*(S_{t+1}, A_{t+1})].$



# Derive TD Target

**Identity:**  $Q^*(s_t, a_t) = \mathbb{E}[R_t + \gamma \cdot Q^*(S_{t+1}, A_{t+1})]$ .

- The action  $A_{t+1}$  is computed by

$$A_{t+1} = \underset{a}{\operatorname{argmax}} Q^*(S_{t+1}, a).$$

- Thus  $Q^*(S_{t+1}, A_{t+1}) = \max_a Q^*(S_{t+1}, a)$ .

# Derive TD Target

**Identity:**  $Q^*(s_t, a_t) = \mathbb{E}[R_t + \gamma \cdot Q^*(S_{t+1}, A_{t+1})]$ .

- Thus  $Q^*(S_{t+1}, A_{t+1}) = \max_a Q^*(S_{t+1}, a)$ .

# Derive TD Target

**Identity:**  $Q^*(s_t, a_t) = \mathbb{E}[R_t + \gamma \cdot Q^*(S_{t+1}, A_{t+1})]$ .

$$= \max_a Q^*(S_{t+1}, a)$$



**Identity:**  $Q^*(s_t, a_t) = \mathbb{E}\left[R_t + \gamma \cdot \max_a Q^*(S_{t+1}, a)\right]$ .

# Derive TD Target

**Identity:**  $Q^*(s_t, a_t) = \mathbb{E} \left[ R_t + \gamma \cdot \max_a Q^*(S_{t+1}, a) \right].$

# Derive TD Target

**Identity:**  $Q^*(s_t, a_t) = \mathbb{E} \left[ \underbrace{R_t}_{\approx r_t} + \gamma \cdot \max_a Q^*(\underbrace{s_{t+1}}_{\Rightarrow s_t}, a) \right].$

$\approx r_t$

# Derive TD Target

**Identity:**  $Q^*(s_t, a_t) = \mathbb{E} \left[ \underbrace{R_t}_{\approx r_t} + \gamma \cdot \underbrace{\max_a Q^*(s_{t+1}, a)}_{\approx \max_a Q^*(s_{t+1}, a)} \right].$



```
graph TD; A["Identity: Q*(s_t, a_t) = E [ R_t + gamma * max_a Q*(s_{t+1}, a) ]"] --> B["≈ r_t"]; A --> C["≈ max_a Q*(s_{t+1}, a)"]
```

$$\approx r_t$$

$$\approx \max_a Q^*(s_{t+1}, a)$$

# Derive TD Target

**Identity:**  $Q^*(s_t, a_t) = \mathbb{E} \left[ \underbrace{R_t + \gamma \cdot \max_a Q^*(s_{t+1}, a)} \right].$



$$\approx r_t + \gamma \cdot \max_a Q^*(s_{t+1}, a)$$

TD target  $y_t$

# **Q-Learning: Tabular Version**



# Q-Learning (tabular version)

- Observe a transition  $(s_t, a_t, r_t, s_{t+1})$ .
- TD target:  $y_t = r_t + \gamma \cdot \max_a Q^*(s_{t+1}, a)$ .

# Q-Learning (tabular version)

- Observe a transition  $(s_t, a_t, r_t, s_{t+1})$ .
- TD target:  $y_t = r_t + \gamma \max_a Q^*(s_{t+1}, a)$ .

	Action $a_1$	Action $a_2$	Action $a_3$	Action $a_4$	...
State $s_1$					
State $s_2$					
State $s_3$					
⋮					

# Q-Learning (tabular version)

- Observe a transition  $(s_t, a_t, r_t, s_{t+1})$ .
- TD target:  $y_t = r_t + \gamma \cdot \max_a Q^*(s_{t+1}, a)$ .
- TD error:  $\delta_t = Q^*(s_t, a_t) - y_t$ .

# Q-Learning (tabular version)

- Observe a transition  $(s_t, a_t, r_t, s_{t+1})$ .
- TD target:  $y_t = r_t + \gamma \cdot \max_a Q^*(s_{t+1}, a)$ .
- TD error:  $\delta_t = Q^*(s_t, a_t) - y_t$ .
- Update:  $Q^*(s_t, a_t) \leftarrow Q^*(s_t, a_t) - \alpha \cdot \delta_t$ .


# **Q-Learning: DQN Version**

# DQN Version

- Approximate  $Q^*(s, a)$  by DQN,  $Q(s, a; \mathbf{w})$ .



state  $s$



DQN

(parameterized by  $\mathbf{w}$ )



$Q(s, \text{"left"}; \mathbf{w})$

$Q(s, \text{"right"}; \mathbf{w})$

$Q(s, \text{"up"}; \mathbf{w})$

# DQN Version

- Approximate  $Q^*(s, a)$  by DQN,  $Q(s, a; \mathbf{w})$ .
- DQN controls the agent by:  $a_t = \underset{a}{\operatorname{argmax}} Q(s_t, a; \mathbf{w})$ .
- We seek to learn the parameter,  $\mathbf{w}$ .

# Q-Learning (DQN Version)

- Observe a transition  $(s_t, a_t, r_t, s_{t+1})$ .
- TD target:  $y_t = r_t + \gamma \cdot \max_a Q(s_{t+1}, a; \mathbf{w})$ .



# Q-Learning (DQN Version)

- Observe a transition  $(s_t, a_t, r_t, s_{t+1})$ .
- TD target:  $y_t = r_t + \gamma \cdot \max_a Q(s_{t+1}, a; \mathbf{w})$ .
- TD error:  $\delta_t = Q(s_t, a_t; \mathbf{w}) - y_t$ .
- Update:  $\mathbf{w} \leftarrow \mathbf{w} - \alpha \cdot \delta_t \cdot \frac{\partial Q(s_t, a_t; \mathbf{w})}{\partial \mathbf{w}}$ .

# Summary

- **Goal:** Learn the optimal action-value function  $Q^*$ .
- **Tabular version** (directly learn  $Q^*$ ).
  - There are finite states and actions.
  - Draw a table, and update the table by Q-learning.
- **DQN version** (function approximation).
  - Approximate  $Q^*$  by the DQN,  $Q(\textcolor{green}{s}, \textcolor{red}{a}; \mathbf{w})$ .
  - Update the parameter,  $\mathbf{w}$ , by Q-learning.

**Thank you!**

<http://wangshusen.github.io/>